

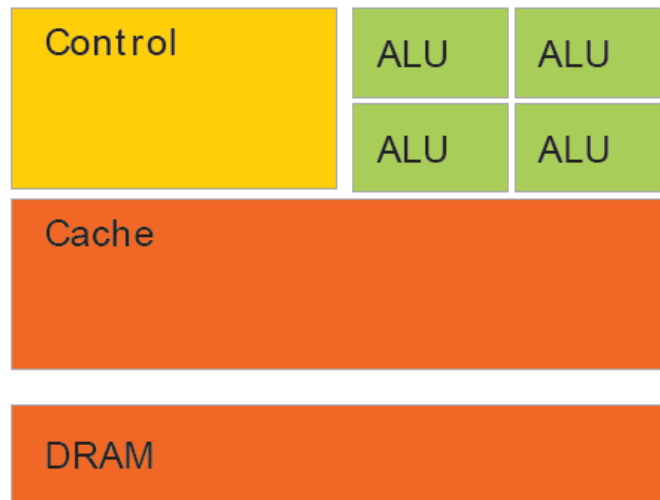
Basic CUDA Programming

Computer Architecture 2020 (Prof. Chih-Wei Liu)

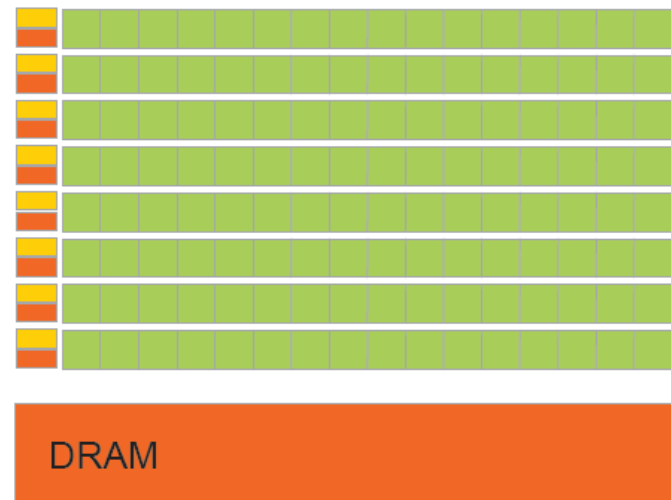
CUDA Tutorial

From Graphics to General Purpose Processing – CPU vs GPU

- ▶ CPU : general purpose computation (SISD)
- ▶ GPU : data-parallel computation (SIMD)



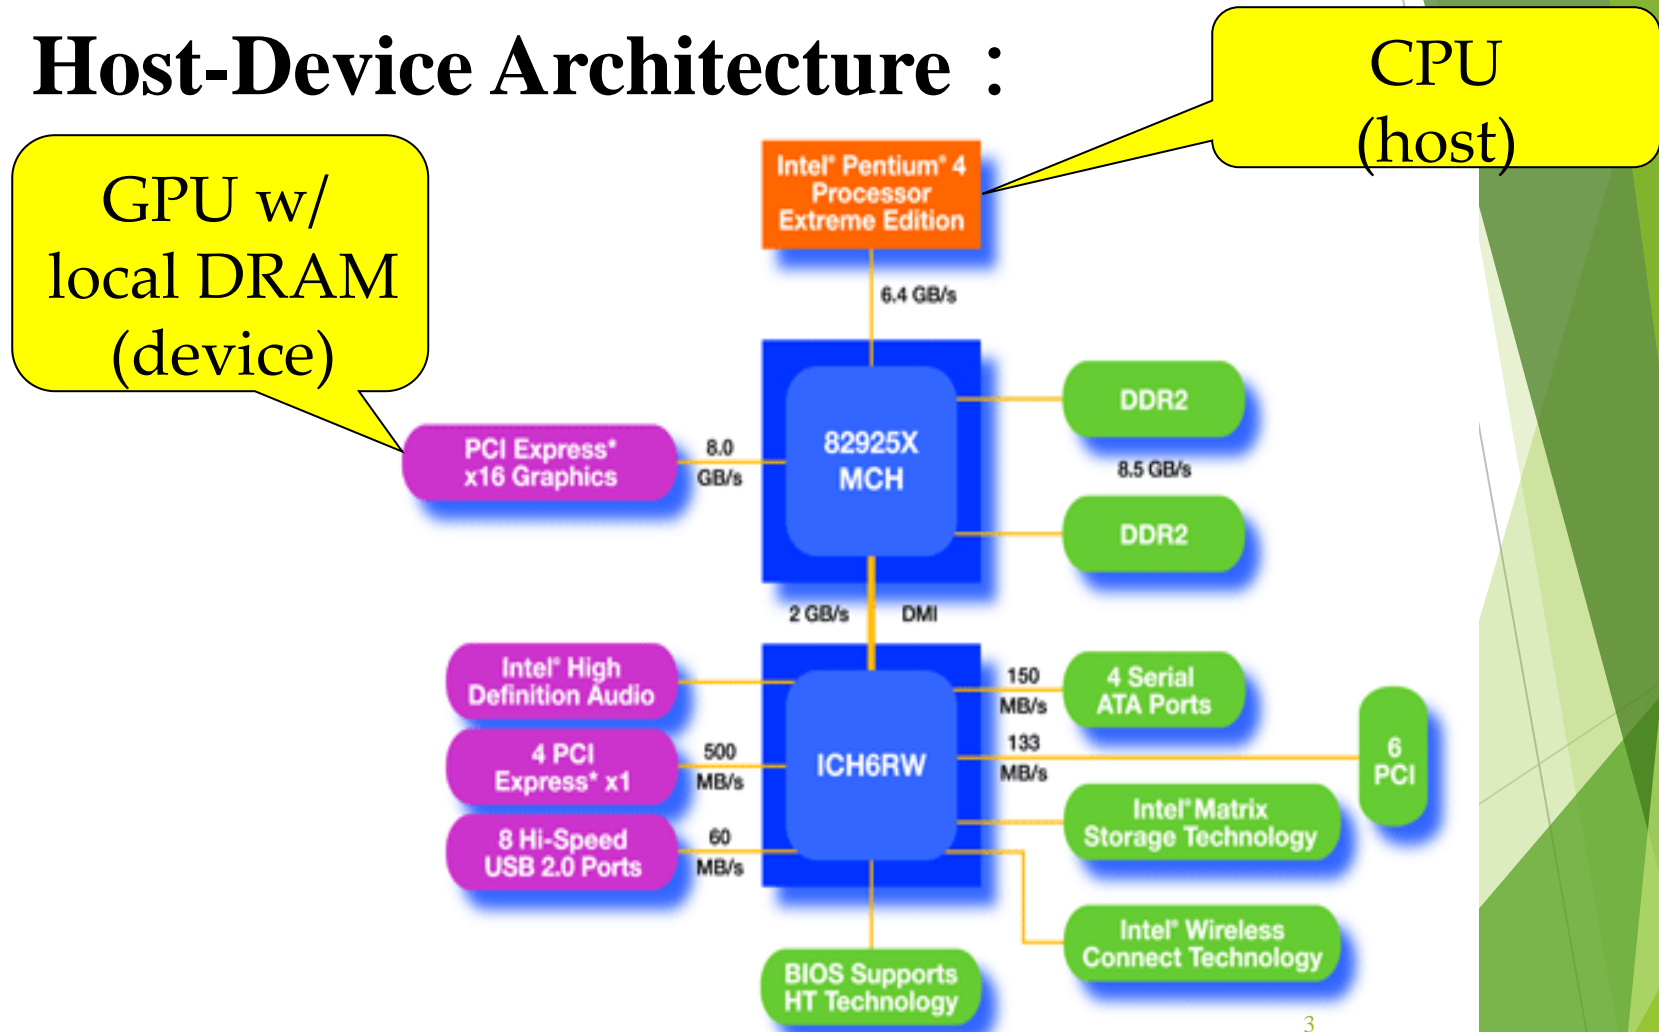
CPU



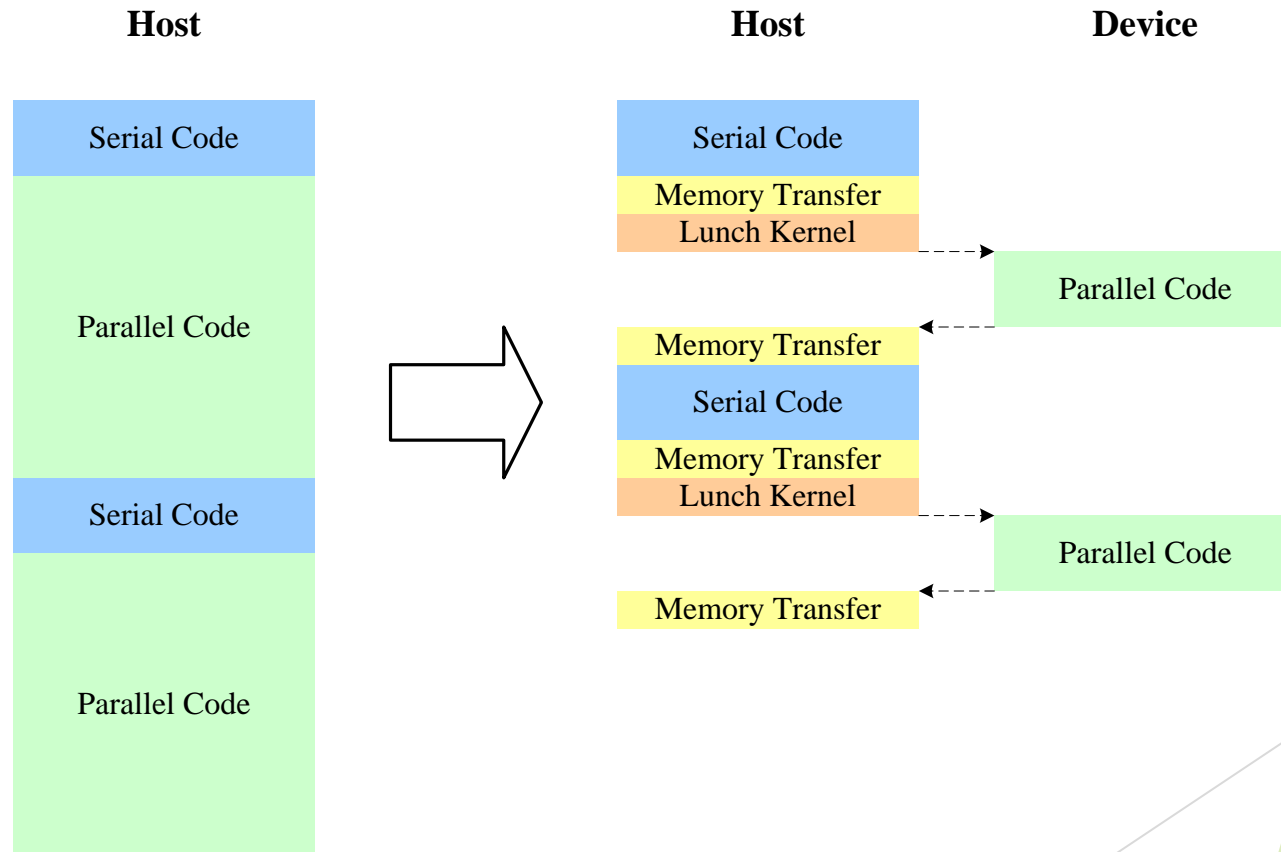
GPU

Heterogeneous computing: CPU+GPU Cooperation

Host-Device Architecture :

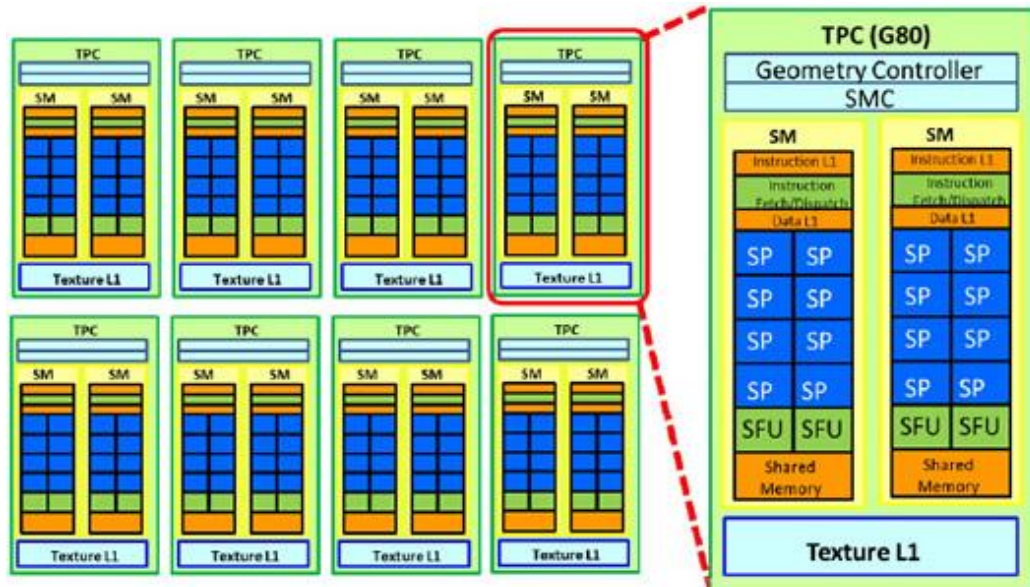


Heterogeneous computing: CUDA Code Execution



Heterogeneous computing: NVIDIA G80 series

- ▶ Streaming Multiprocessor (SM)
 - ▶ Streaming Processor (SP) / [CUDA core](#)
 - ▶ Shared memory
- ▶ For NV GTX 1080, the number of SPs is 2560.



CUDA Programming Model (1/7)

- ▶ Define
 - ▶ Programming model
 - ▶ Memory model
- ▶ Help developers map the applications or algorithms onto CUDA devices more easily and clearly.
- ▶ It is important to follow CUDA's programming model to obtain higher performance of program execution.

CUDA Programming Model (2/7)

► C/C++ for CUDA

- Subset of C with extensions
- C++ templates for GPU code
- CUDA goals:
 - Scale code to 100s of cores and 1000s of parallel threads.
 - Facilitate heterogeneous computing: CPU + GPU

CUDA Programming Model (3/7)

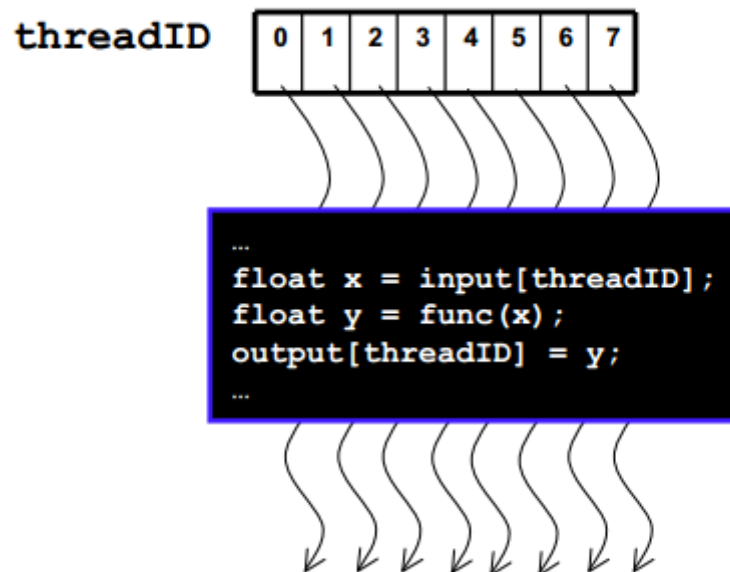
▶ CUDA Kernels and Threads :

- ▶ Parallel portions of an application executed on the device are **kernels**. And only one **kernel** is executed at a time.
- ▶ All the **threads** execute the same kernel at a time.

CUDA Programming Model (4/7)

Arrays of Parallel Threads :

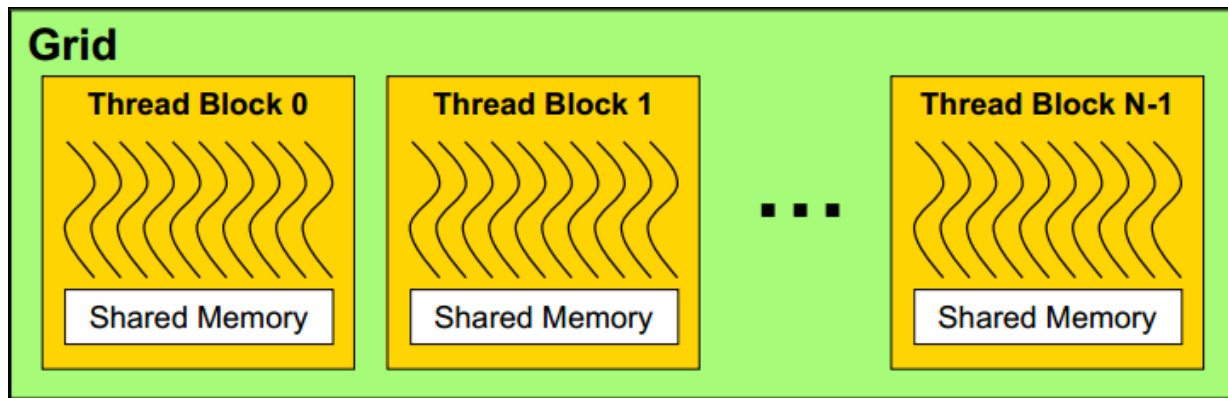
- ▶ A CUDA kernel is executed by an array of threads
 - ▶ Each thread has an ID that it uses to compute memory addresses and make control decisions



CUDA Programming Model (5/7)

Thread Batching :

- ▶ Kernel launches a **grid** of **thread blocks**

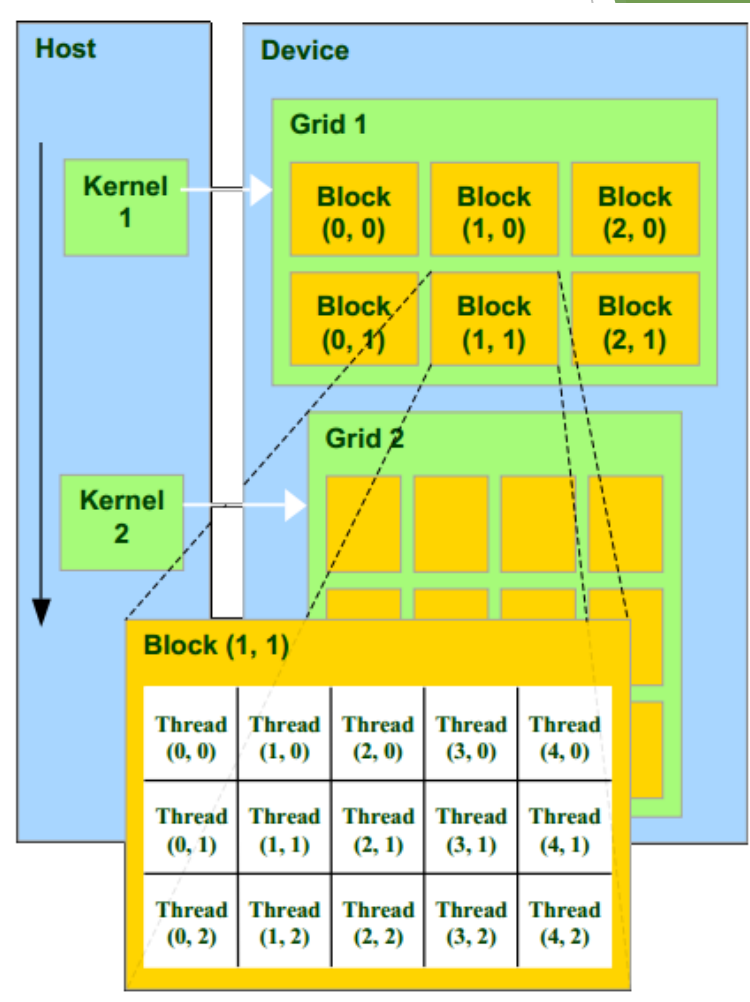


- ▶ Threads within a block cooperate via shared memory
- ▶ Threads in different blocks cannot cooperate

CUDA Programming Model (6/7)

CUDA Programming Model :

- ▶ A kernel is executed by a grid of thread blocks
 - ▶ Block can be 1D or 2D or 3D.
- ▶ A thread block is a batch of threads
 - ▶ Thread can be 1D or 2D or 3D.
 - ▶ Data can be shared through shared memory
 - ▶ Execution synchronization
 - ▶ But threads from different blocks can't cooperate.



CUDA Programming Model (7/7)

Memory Model :

▶ Registers

- ▶ Per thread
- ▶ Data lifetime = thread lifetime

▶ Shared memory

- ▶ Per thread block on-chip memory
- ▶ Data lifetime = block lifetime

▶ Local memory

- ▶ Per thread off-chip memory (physically in device DRAM)
- ▶ Data lifetime = thread lifetime

▶ Global (device) memory

- ▶ Accessible by all threads as well as host (CPU)
- ▶ Data lifetime = from allocation to deallocation

▶ Host (CPU) memory

- ▶ Not directly accessible by CUDA threads

Heterogeneous computing: NVIDIA CUDA Compiler (NVCC)

- ▶ NVCC separates CPU and GPU source code into two parts.
 - ▶ For host codes, NVCC invokes typical C compiler like GCC, Intel C compiler, or MS C compiler.
 - ▶ All the device codes are compiled by NVCC.
 - ▶ The extension of device source files should be “.cu”.
- ▶ All executable with CUDA code requires :
 - ▶ CUDA core library (cuda)
 - ▶ CUDA runtime library (cudart)

CUDA C/C++ Basic

GPU Memory Allocation/Release

- ▶ Memory allocation on GPU
 - ▶ `cudaMalloc(void **pointer, size_t nbytes)`
- ▶ Preset value for specific memory area
 - ▶ `cudaMemset(void *pointer, int value, size_t count)`
- ▶ Release memory allocation
 - ▶ `cudaFree(void *pointer)`

```
int n = 1024;  
int nbytes = 1024*sizeof(int);  
int *d_a = 0;  
cudaMalloc( (void**)&d_a, nbytes );  
cudaMemset( d_a, 0, nbytes);  
cudaFree(d_a);
```

Data Copies

- ▶ `cudaMemcpy(void *dst, void *src, size_t nbytes, enum cudaMemcpyKind direction);`
 - ▶ *direction* specifies locations (host or device) of *src* and *dst*
 - ▶ Blocks CPU thread: returns after the copy is complete
 - ▶ Doesn't start copying until previous CUDA calls complete
- ▶ `enum cudaMemcpyKind`
 - ▶ `cudaMemcpyHostToDevice`
 - ▶ `cudaMemcpyDeviceToHost`
 - ▶ `cudaMemcpyDeviceToDevice`

Function Qualifiers

- ▶ `__global__` : invoked from within host (CPU) code, also called 'kernel'
- ▶ `__device__` : called from other GPU functions, cannot be called from host (CPU) code
- ▶ `__host__` : can only be executed by CPU, called from host

Variable Qualifiers (GPU code)

▶ `__device__`

- ▶ Stored in device memory (large capacity, high latency, uncached)
- ▶ Allocated with `cudaMalloc` (`__device__` qualifier implied)
- ▶ Accessible by all threads
- ▶ Lifetime: application

▶ `__shared__`

- ▶ Stored in on-chip shared memory (SRAM, low latency)
- ▶ Allocated by execution configuration or at compile time
- ▶ Accessible by all threads in the same thread block
- ▶ Lifetime: duration of thread block

CUDA Built-in Device Variables

- ▶ All `__global__` and `__device__` functions have access to these automatically defined variables
- ▶ `dim3 gridDim;`
 - ▶ Dimensions of the grid in blocks
- ▶ `dim3 blockIdx;`
 - ▶ Block index within the grid
- ▶ `dim3 blockDim;`
 - ▶ Dimensions of the block in threads
- ▶ `dim3 threadIdx;`
 - ▶ Thread index within the block

Launching Kernels

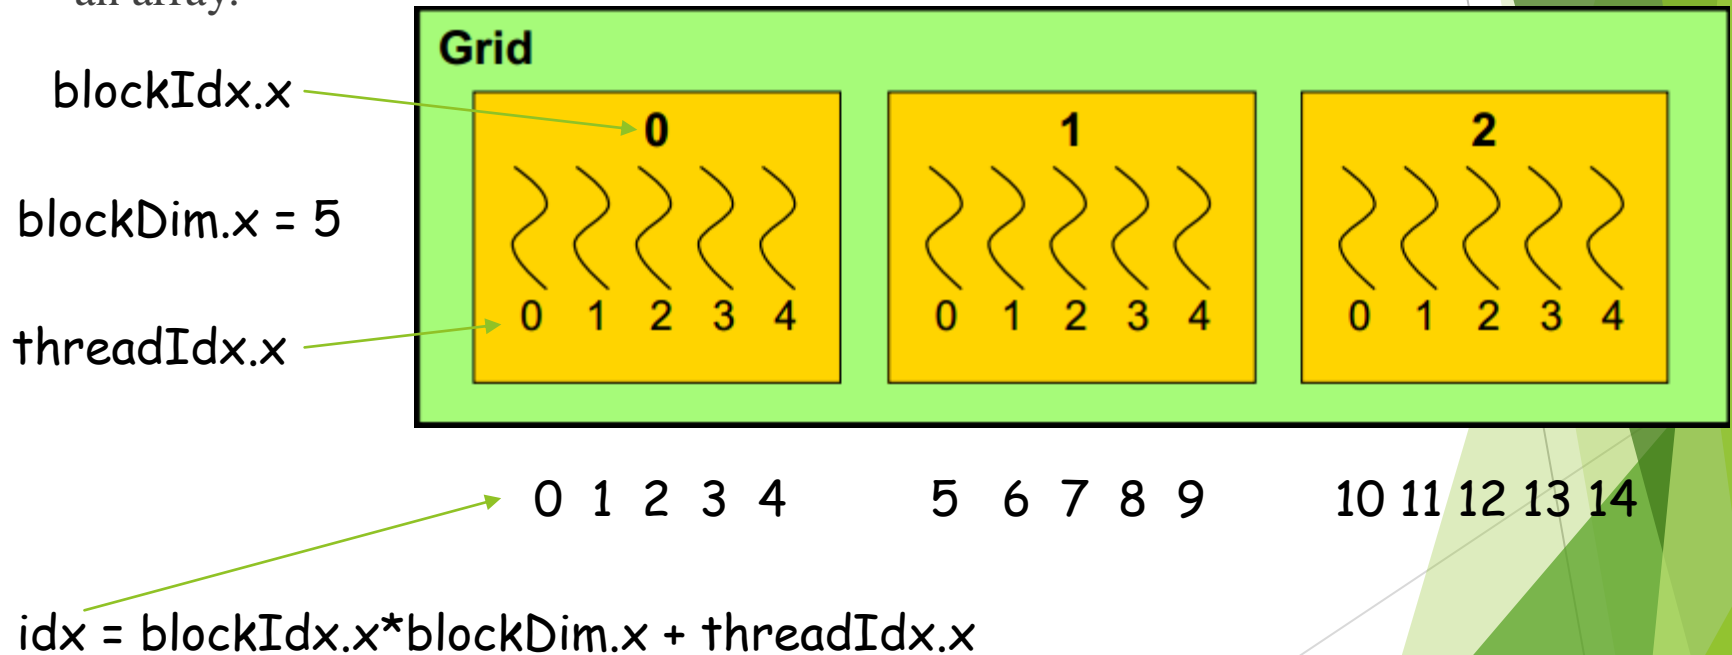
► `kernel<<<dim3 grid, dim3 block>>>(...);`

```
dim3 grid(16,16);  
dim3 block(16,16);  
kernel1<<<grid, block>>>(...);
```

```
kernel2<<<32, 512>>>(...);
```

Data Decomposition

- Often want each thread in kernel to access a different element of an array.



Data Decomposition Example: Increment Array Elements (1) (1/4)

► Increment N-element vector **a** by scalar **b**

► Each thread only executes ONCE

CPU program

```
void increment_cpu(int a[], int b, int N)
{
    for(int idx=0;idx<N;idx++)
        a[idx]=a[idx]+b;
}
```

```
void main()
{
    ...
    increment_cpu(a,b,N);
}
```

CUDA program

```
__global__ void increment_gpu(int a[], int b, int N)
{
    int idx = blockIdx.x*blockDim.x+threadIdx.x;
    if(idx<N)
        a[idx]=a[idx]+b;
}
```

```
void main()
{
    ...
    dim3 dimBlock(BlockSize);
    dim3 dimGrid(ThreadNum);
    increment_gpu<<<dimGrid,dimBlock>>>(d_a,b,N);
}
```



Data Decomposition Example: Increment Array Elements (1) (2/4)

```
#include <iostream>
#include <cuda.h>
#include <cuda_runtime.h>
using namespace std;

const int BlockSize = 2;
const int ThreadNum = 5;
const int N = 10;

__global__ void increment_gpu(int a[], const int b, const int N)
{
    int idx = blockIdx.x*blockDim.x+threadIdx.x;
    if(idx<N)
        a[idx]=a[idx]+b;
}
```

1. Allocate an integer array on device's memory
2. Copy the 'a' array's value to 'd_a' (on the device)

```
int main()
{
    int a[N];
    int *d_a = 0;
    int i = 0;
    const int b = 10;

    cout << "a[N] array before scaling: [";
    for(i=0;i<N;i++)
    {
        a[i] = i;
        cout << a[i] << " ";
    }
    cout << "]" << endl;

    cudaMalloc((void**) &d_a, sizeof(int) * N);
    cudaMemcpy(d_a, a, sizeof(int) * N, cudaMemcpyHostToDevice);

    dim3 dimBlock(BlockSize);
    dim3 dimGrid(ThreadNum);
    increment_gpu<<<dimGrid,dimBlock>>>>(d_a,b,N);
    cudaDeviceSynchronize();

    cudaMemcpy(a,d_a,sizeof(int) * N,cudaMemcpyDeviceToHost);

    cout << "a[N] array after scaling: [";
    for(i=0;i<N;i++)
    {
        cout << a[i] << " ";
    }
    cout << "]" << endl;

    return 0;
}
```

Pointer referred to device memory space

Data Decomposition Example: Increment Array Elements (1) (3/4)

```
#include <iostream>
#include <cuda.h>
#include <cuda_runtime.h>
using namespace std;

const int BlockSize = 2;
const int ThreadNum = 5;
const int N = 10;

__global__ void increment_gpu(int a[],const int b, const int N)
{
    int idx = blockIdx.x*blockDim.x+threadIdx.x;
    if(idx<N)
        a[idx]=a[idx]+b;
}
```

3. Invoke CUDA kernel function

4. Make sure all the tasks are finished

5. Retrieve the results from device.

main_gpu_1.cu

```
int main()
{
    int a[N];
    int *d_a = 0;
    int i = 0;
    const int b = 10;

    cout << "a[N] array before scaling: [";
    for(i=0;i<N;i++)
    {
        a[i] = i;
        cout << a[i] << " ";
    }
    cout << "]" << endl;

    cudaMalloc((void**) &d_a,sizeof(int)*N);
    cudaMemcpy(d_a,a,sizeof(int)*N,cudaMemcpyHostToDevice);

    dim3 dimBlock(BlockSize);
    dim3 dimGrid(ThreadNum);
    increment_gpu<<<dimGrid,dimBlock>>>(d_a,b,N);
    cudaDeviceSynchronize();

    cudaMemcpy(a,d_a,sizeof(int)*N,cudaMemcpyDeviceToHost);

    cout << "a[N] array after scaling: [";
    for(i=0;i<N;i++)
    {
        cout << a[i] << " ";
    }
    cout << "]" << endl;

    return 0;
}
```


Data Decomposition Example: Increment Array Elements (1) (4/4)

$N = 10$ \leftarrow total tasks

BlockSize = 2

ThreadNum = 5 \leftarrow total threads = $2 * 5$

	a[0]	a[1]	a[2]	a[3]	a[4]	a[5]	a[6]	a[7]	a[8]	a[9]
#1	(0,0)	(0,1)	(0,2)	(0,3)	(0,4)	(1,0)	(1,1)	(1,2)	(1,3)	(1,4)

(blockIdx.x , threadIdx.x)

Data Decomposition Example: Increment Array Elements (2) (1/3)

$N = 10$ ← total tasks

BlockSize = 1

ThreadNum = 5 ← total threads = $1 * 5$

	a[0]	a[1]	a[2]	a[3]	a[4]	a[5]	a[6]	a[7]	a[8]	a[9]
#1	(0,0)		(0,1)		(0,2)		(0,3)		(0,4)	
#2		(0,0)		(0,1)		(0,2)		(0,3)		(0,4)

← stripe size →

(blockIdx.x , threadIdx.x)

Data Decomposition Example: Increment Array Elements (2) (2/3)

► Increment N-element vector **a** by scalar **b**

CPU program

```
void increment_cpu(float *a, float *b, int N)
{
    for(int idx=0;idx<N;idx++)
        a[idx]=a[idx]+b;
}

void main()
{
    ...
    increment_cpu(a,b,N);
}
```

CUDA program

```
__global__ void increment_gpu(float *a, float *b, int N)
{
    int TotalThread = blockDim.x*blockDim.x;
    int stripe = N / TotalThread;
    int head = (blockIdx.x*blockDim.x + threadIdx.x)*stripe;
    for(int i = head;i<(head+stripe);i++)
        a[i]=a[i]+b;
}

void main()
{
    ...
    dim3 dimBlock(blocksize);
    dim3 dimGrid(ceil(N/(float)blocksize));
    increment_gpu<<<dimGrid,dimBlock>>>>(a,b,N);
}
```



Data Decomposition Example: Increment Array Elements (2) (3/3)

```
#include <iostream>
#include <cuda.h>
#include <cuda_runtime.h>
using namespace std;

const int BlockSize = 1;
const int ThreadNum = 5;
const int N = 10;

__global__ void increment_gpu(int a[],const int b, const int N)
{
    int i = 0;
    int TotalThread = gridDim.x*blockDim.x;
    int stripe = N / TotalThread;
    int head = (blockIdx.x*blockDim.x + threadIdx.x)*stripe;

    for(i = head;i<(head+stripe);i++)
        a[i]=a[i]+b;
}
```

```
int main()
{
    int a[N];
    int *d_a = 0;
    int i = 0;
    const int b = 10;

    cout << "a[N] array before scaling: [";
    for(i=0;i<N;i++)
    {
        a[i] = i;
        cout << a[i] << " ";
    }
    cout << "]" << endl;

    cudaMalloc((void**) &d_a,sizeof(int)*N);
    cudaMemcpy(d_a,a,sizeof(int)*N,cudaMemcpyHostToDevice);

    dim3 dimBlock(BlockSize);
    dim3 dimGrid(ThreadNum);
    increment_gpu<<<dimGrid,dimBlock>>>>(d_a,b,N);
    cudaDeviceSynchronize();

    cudaMemcpy(a,d_a,sizeof(int)*N,cudaMemcpyDeviceToHost);

    cout << "a[N] array after scaling: [";
    for(i=0;i<N;i++)
    {
        cout << a[i] << " ";
    }
    cout << "]" << endl;

    return 0;
}
```

Reference

- ▶ NVIDIA CUDA Toolkit Documentations
(<http://docs.nvidia.com/cuda/>)
- ▶ Recommended learning materials
 - ▶ CUDA Programming C/C++ basic
(https://www.youtube.com/watch?v=kyL2rj_Se3M)